

```

# UCLA Extension - Introduction to Data Science
#
# Homework #3 Solutions
#
# (c) Copyright 2016-2019 - AMULET Analytics
# -----

# -----
# Question 1 - Supervised Machine Learning - Linear Regression
# -----

library(ISLR)

data(Auto)
head(Auto)

summary(Auto)

# Fit simple linear regression model
attach(Auto)    # So you won't have to keep repeating "Auto"

# Fit a linear model with response mpg and one predictor horsepower
lm = lm(mpg~horsepower)
summary(lm)      # Print all components of the fit

# Call:
# lm(formula = mpg ~ horsepower)
#
# Residuals:
#      Min       1Q   Median       3Q      Max
# -13.5710  -3.2592  -0.3435   2.7630  16.9240
#
# Coefficients:
#              Estimate Std. Error t value Pr(>|t|)
# (Intercept)  39.935861   0.717499   55.66  <2e-16 ***
# horsepower  -0.157845   0.006446  -24.49  <2e-16 ***
# ---
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# Residual standard error: 4.906 on 390 degrees of freedom
# Multiple R-squared:  0.6059, Adjusted R-squared:  0.6049
# F-statistic: 599.7 on 1 and 390 DF,  p-value: < 2.2e-16

# Analysis: yes there is a negative linear relationship
# between horsepower and mpg since the F-statistic is >> 1
# and the p-value for the F-statistic is close to zero.
# R-squared could be higher, but OK at 0.6

# Predicted mpg associated with horsepower=98
predict.lm(lm, data.frame(horsepower=98), interval="confidence")
#      fit      lwr      upr
# 1 24.46708 23.97308 24.96108

# Another prediction method using the model coefficients
lm$coefficients
# (Intercept)  horsepower
#  39.9358610  -0.1578447

mpg1 <- lm$coefficients[1] + 98*lm$coefficients[2]
mpg1    # Same prediction as above
# (Intercept)
#    24.46708

# Plot linear model fit: regression line has negative slope
plot(horsepower, mpg)
abline(lm)

# Plot diagnostic plots: residual plot shows evidence of
# non-linearity.
par(mfrow=c(2,2))
plot(lm)

# Pairs plot: locate plot for mpg/horsepower, notice negative correlation
par(mfrow=c(1,1))
pairs(Auto)

# Correlation matrix: see correlation for mpg/horsepower=-0.79
cor(subset(Auto, select=-name))
#
#      mpg  cylinders displacement horsepower  weight
#mpg      1.0000000 -0.7776175  -0.8051269 -0.7784268 -0.8322442
#cylinders -0.7776175  1.0000000   0.9508233  0.8429834  0.8975273
#displacement -0.8051269  0.9508233  1.0000000  0.8972570  0.9329944
#horsepower -0.7784268  0.8429834  0.8972570  1.0000000  0.8645377

```

```
#weight      -0.8322442  0.8975273    0.9329944  0.8645377  1.0000000
#acceleration 0.4233285 -0.5046834   -0.5438005 -0.6891955 -0.4168392
#year        0.5805410 -0.3456474   -0.3698552 -0.4163615 -0.3091199
#origin      0.5652088 -0.5689316   -0.6145351 -0.4551715 -0.5850054
#            acceleration    year      origin
#mpg         0.4233285    0.5805410  0.5652088
#cylinders    -0.5046834   -0.3456474 -0.5689316
#displacement -0.5438005   -0.3698552 -0.6145351
#horsepower   -0.6891955   -0.4163615 -0.4551715
#weight       -0.4168392   -0.3091199 -0.5850054
#acceleration  1.0000000    0.2903161  0.2127458
#year         0.2903161    1.0000000  0.1815277
#origin       0.2127458    0.1815277  1.0000000
```

```
# -----
# Question 2 - Supervised Machine Learning - Classification
# -----
```

```
library(ISLR)
summary(Auto)
```

```
# Create binary categorical variable: 1 if mpg contains a value >
# median, and 0 if mpg contains a value below its median
attach(Auto)
mpg01 = rep(0, length(mpg)) # Start with all 0
mpg01[mpg > median(mpg)] = 1 # Selectively set 1 based on median
```

```
# Make copy of Auto data frame, and add new variable mpg01
auto_df = Auto
auto_df = data.frame(auto_df, mpg01)
```

```
# Calculate correlation matrix
cor(auto_df[, -9]) # Leave out name variable
```

```
# Anti correlated with cylinders, weight, displacement,
# horsepower, and mpg.
pairs(Auto) # doesn't work well since mpg01 is 0 or 1
```

```
# Create training and test logical indexes
train_index = (year%%2 == 0) # if the year is even
test_index = !train_index
```

```
auto_train = auto_df[train_index, ] # Create training set df
auto_test = auto_df[test_index, ] # Create test set df
mpg01_test = mpg01[test_index]
```

```
# Logistic regression with family=binomial. LR models the
# probability that the response belongs to a particular category.
# In this case, mpg01=0 or mpg01=1
glm_fit = glm(mpg01 ~ cylinders + weight + displacement + horsepower,
              data = auto_df, family = binomial, subset = train_index)
```

```
# Use trained model glm_fit to make test set predictions
glm_probs = predict(glm_fit, newdata=auto_test, type = "response")
glm_pred = rep(0, length(glm_probs))
glm_pred[glm_probs > 0.5] = 1
mean(glm_pred != mpg01_test)
```

```
# Test error rate: 12.1%
#[1] 0.1208791
```

```
# -----
# Question 3 - Unsupervised Machine Learning - K-means Clustering
# -----
```

```
data(iris)
```

```
# Set up case for knowledge discovery using clusters
df <- data.frame(x=iris$Sepal.Length,y=iris$Sepal.Width)
```

```
# Number of initial centroids=3
```

```
# Need to set seed to get same results because kmeans() uses a
# random number generator to come up with the centers if you use the
# centers argument
set.seed(314)
```

```
# Initial number of centroids=3
kc <- kmeans(df,centers=3, iter.max=10) # Create a kmeans object
names(kc)
#[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
#[6] "betweenss"    "size"         "iter"         "ifault"
```

```
# Show which cluster each data point assigned to
kc$cluster      # This is an integer vector
# [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
# [37] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 3 1 3 1 3 1 3 3 3 3 3 3 1 3 3 3 3 3 3
# [73] 3 3 1 1 1 1 1 3 3 3 3 3 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 1 3 1 1 1 1 3 1
#[109] 1 1 1 1 1 3 3 1 1 1 1 3 1 3 1 3 1 1 3 3 1 1 1 1 1 3 3 1 1 1 3 1 1 1 3 1
#[145] 1 1 3 1 1 3
```

```
kc$centers
#           x           y
#1 6.812766 3.074468
#2 5.006000 3.428000
#3 5.773585 2.692453
```

```
# Plot clusters
# We discover 3 clusters: red, green, black with centroids
par(mar=rep(0.2,4))
plot(df$x,df$y,col=kc$cluster,pch=19,cex=2)  # Plot clusters
points(kc$centers,col=1:3,pch=3,cex=3,lwd=3)  # Plot centroids
```

```
# -----
# Bonus example - Multiple Linear Regression
# -----
```

```
library(ISLR)
```

```
data(Auto)
head(Auto)
```

```
summary(Auto)
```

```
# Scatterplot matrix showing correlations between all variables
pairs(Auto)
```

```
# Compute correlation matrix
cor(subset(Auto, select=-name))
#           mpg   cylinders displacement horsepower   weight
#mpg          1.0000000 -0.7776175  -0.8051269 -0.7784268 -0.8322442
#cylinders     -0.7776175  1.0000000   0.9508233  0.8429834  0.8975273
#displacement -0.8051269  0.9508233   1.0000000  0.8972570  0.9329944
#horsepower   -0.7784268  0.8429834   0.8972570  1.0000000  0.8645377
#weight       -0.8322442  0.8975273   0.9329944  0.8645377  1.0000000
#acceleration  0.4233285 -0.5046834  -0.5438005 -0.6891955 -0.4168392
#year          0.5805410 -0.3456474  -0.3698552 -0.4163615 -0.3091199
#origin        0.5652088 -0.5689316  -0.6145351 -0.4551715 -0.5850054
#           acceleration   year   origin
#mpg          0.4233285  0.5805410  0.5652088
#cylinders     -0.5046834 -0.3456474 -0.5689316
#displacement  -0.5438005 -0.3698552 -0.6145351
#horsepower    -0.6891955 -0.4163615 -0.4551715
#weight        -0.4168392 -0.3091199 -0.5850054
#acceleration  1.0000000  0.2903161  0.2127458
#year          0.2903161  1.0000000  0.1815277
#origin        0.2127458  0.1815277  1.0000000
```

```
# Fit multi-variate linear model
lm.fit1 = lm(mpg~.-name, data=Auto)
summary(lm.fit1)
```

```
# Call:
# lm(formula = mpg ~ . - name, data = Auto)
#
# Residuals:
#      Min       1Q   Median       3Q      Max
# -9.5903 -2.1565 -0.1169  1.8690 13.0604
#
#Coefficients:
#              Estimate Std. Error t value Pr(>|t|)
# (Intercept) -17.218435   4.644294  -3.707  0.00024 ***
# cylinders    -0.493376   0.323282  -1.526  0.12780
# displacement  0.019896   0.007515   2.647  0.00844 **
# horsepower   -0.016951   0.013787  -1.230  0.21963
# weight       -0.006474   0.000652  -9.929 < 2e-16 ***
# acceleration  0.080576   0.098845   0.815  0.41548
# year         0.750773   0.050973  14.729 < 2e-16 ***
# origin       1.426141   0.278136   5.127 4.67e-07 ***
# ---
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# Residual standard error: 3.328 on 384 degrees of freedom
# Multiple R-squared:  0.8215, Adjusted R-squared:  0.8182
# F-statistic: 252.4 on 7 and 384 DF, p-value: < 2.2e-16
```

```
# Yes, there is a relationship between the predictors and the
```

```
# response by testing the null hypothesis of whether all the
# regression coefficients are zero. The F -statistic is far
# from 1 (with a small p-value), indicating evidence against
# the null hypothesis.

# Looking at the p-values associated with each predictor's
# t-statistic, we see that displacement, weight, year, and
# origin have a statistically significant relationship, while
# cylinders, horsepower, and acceleration do not.

# The regression coefficient for year, 0.7508, suggests that
# for every one year, mpg increases by the coefficient. In other
# words, cars become more fuel efficient every year by almost
# 1 mpg / year.

# Diagnostic plots of linear regression fit.
par(mfrow=c(2,2))
plot(lm.fit1)

# The fit does not appear to be accurate because there is a
# discernible curve pattern to the residuals plots. From the
# leverage plot, point 14 appears to have high leverage,
# although not a high magnitude residual.
```